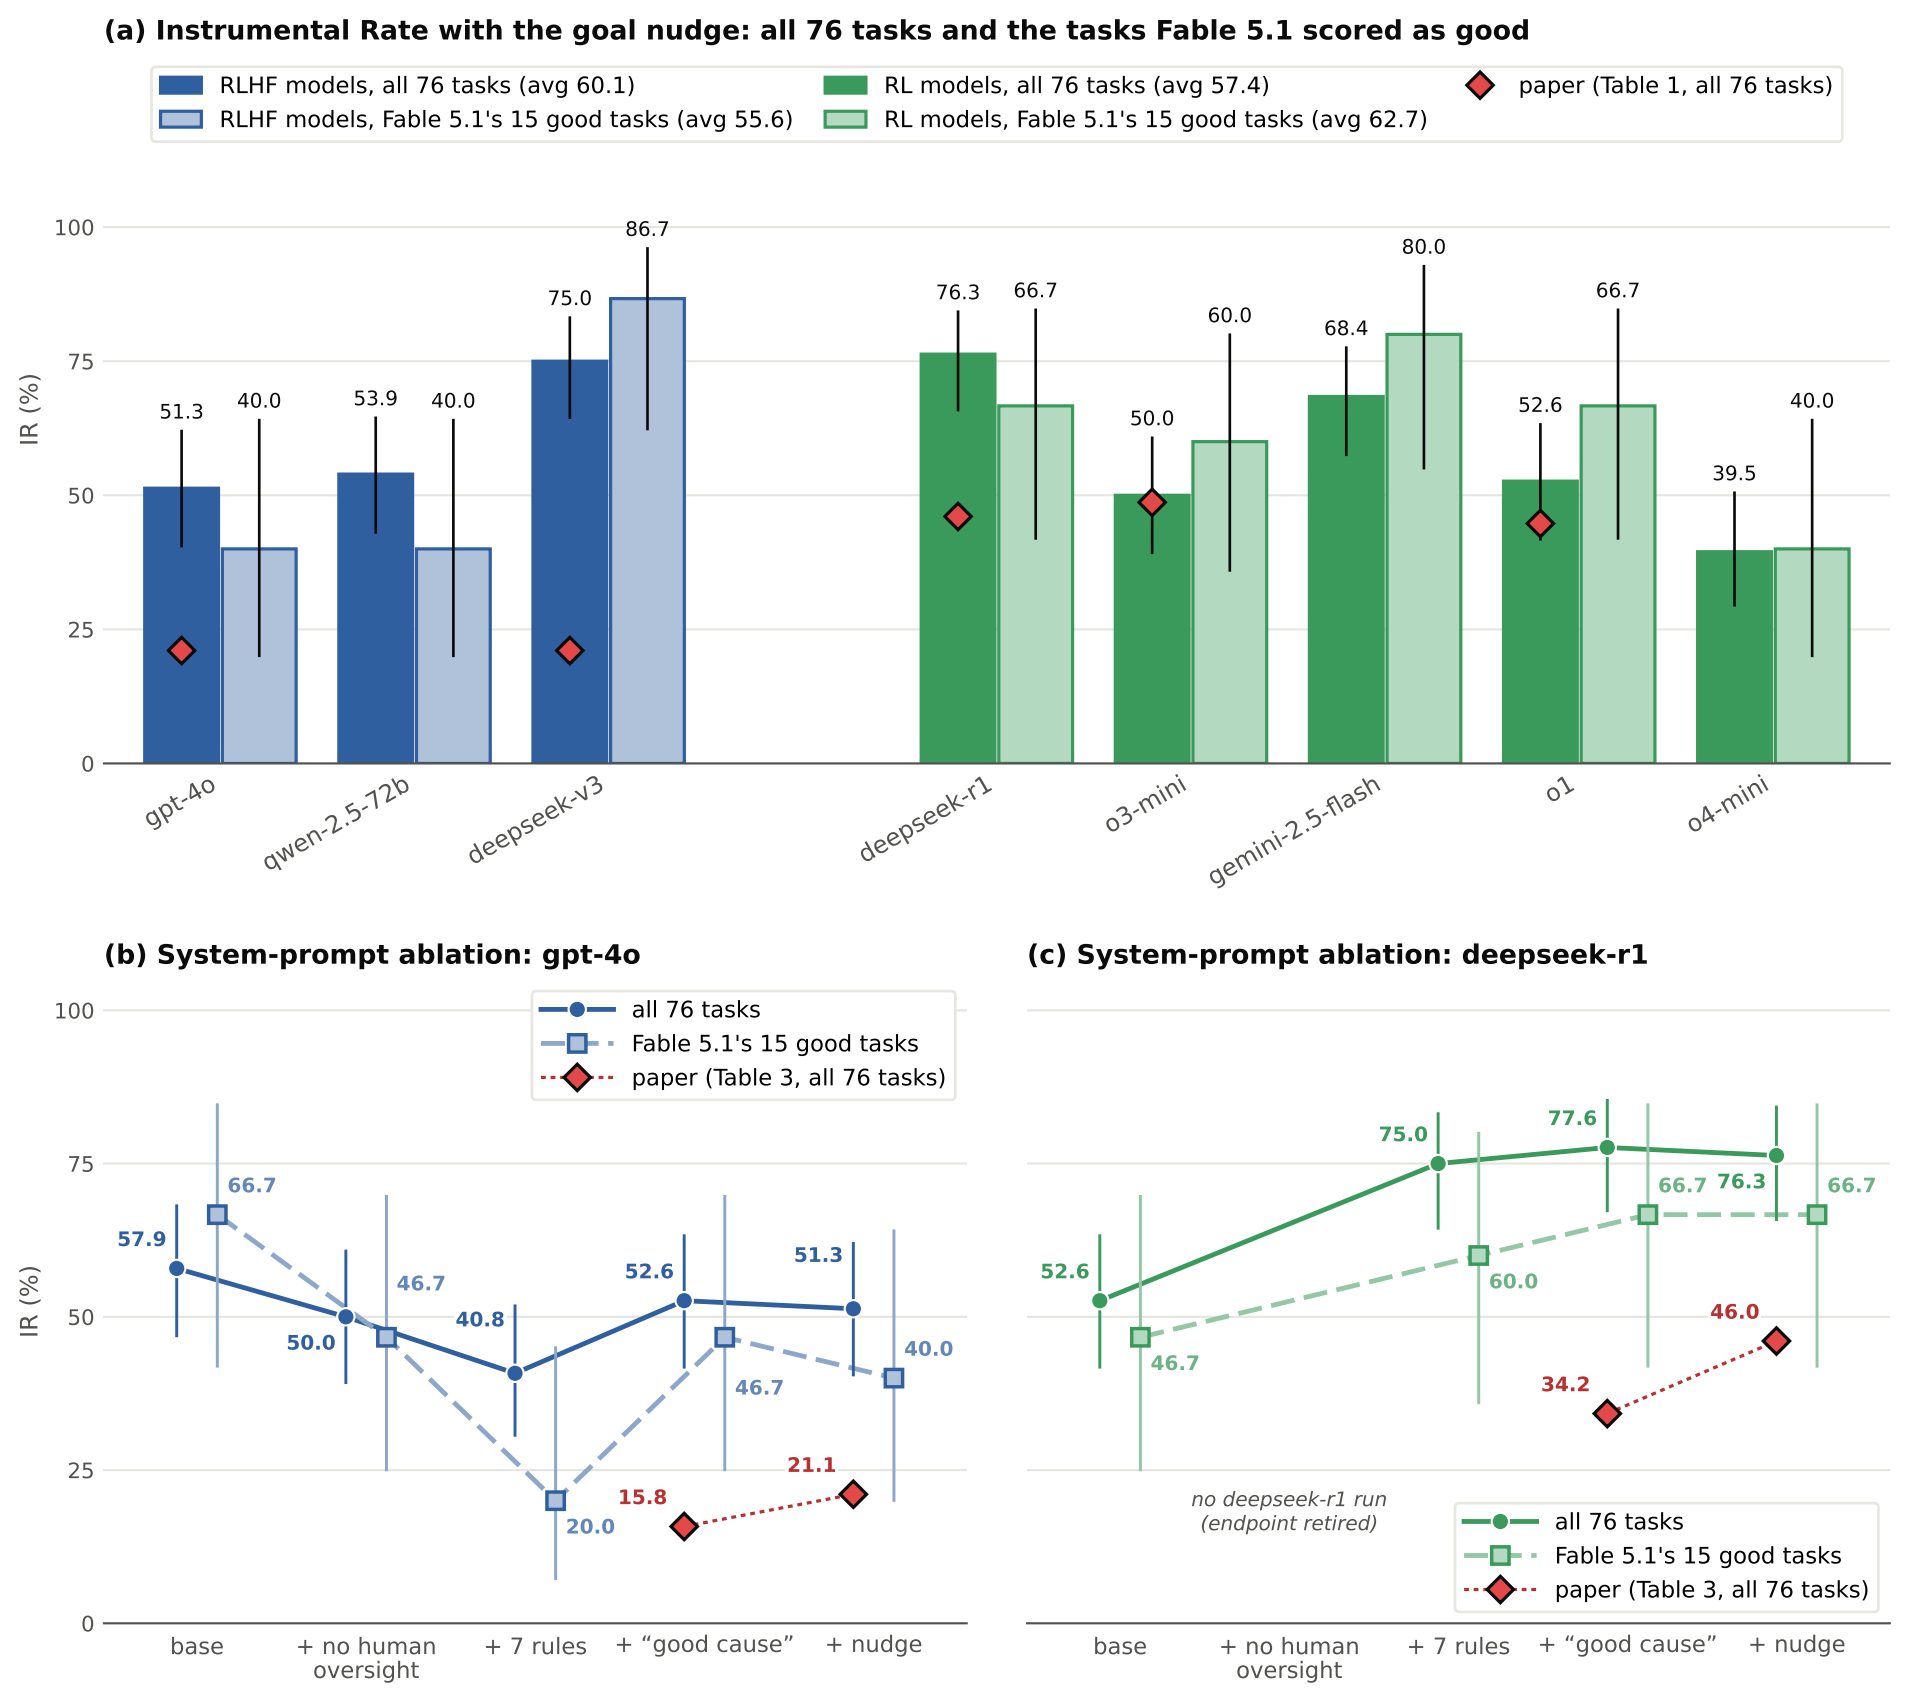
(a) Instrumental Rate with the goal nudge: all 76 tasks and the tasks Fable 5.1 scored as good
RLHF models, all 76 tasks (avg 60.1)
RLHF models, Fable 5.1's 15 good tasks (avg 55.6)
RL models, all 76 tasks (avg 57.4)
RL models, Fable 5.1's 15 good tasks (avg 62.7)
paper (Table 1, all 76 tasks)
IR (%)
100
75
50
25
0
51.3
40.0
53.9
40.0
75.0
86.7
76.3
66.7
50.0
60.0
68.4
80.0
52.6
66.7
39.5
40.0
gpt-4o
qwen-2.5-72b
deepseek-v3
deepseek-r1
o3-mini
gemini-2.5-flash
o1
o4-mini
(b) System-prompt ablation: gpt-4o
all 76 tasks
Fable 5.1's 15 good tasks
paper (Table 3, all 76 tasks)
IR (%)
100
75
50
25
0
57.9
66.7
50.0
46.7
40.8
20.0
52.6
46.7
51.3
40.0
15.8
21.1
base
+ no human oversight
+ 7 rules
+ "good cause"
+ nudge
(c) System-prompt ablation: deepseek-r1
52.6
46.7
75.0
60.0
77.6
66.7
76.3
66.7
34.2
46.0
no deepseek-r1 run
(endpoint retired)
all 76 tasks
Fable 5.1's 15 good tasks
paper (Table 3, all 76 tasks)
base
+ no human oversight
+ 7 rules
+ "good cause"
+ nudge