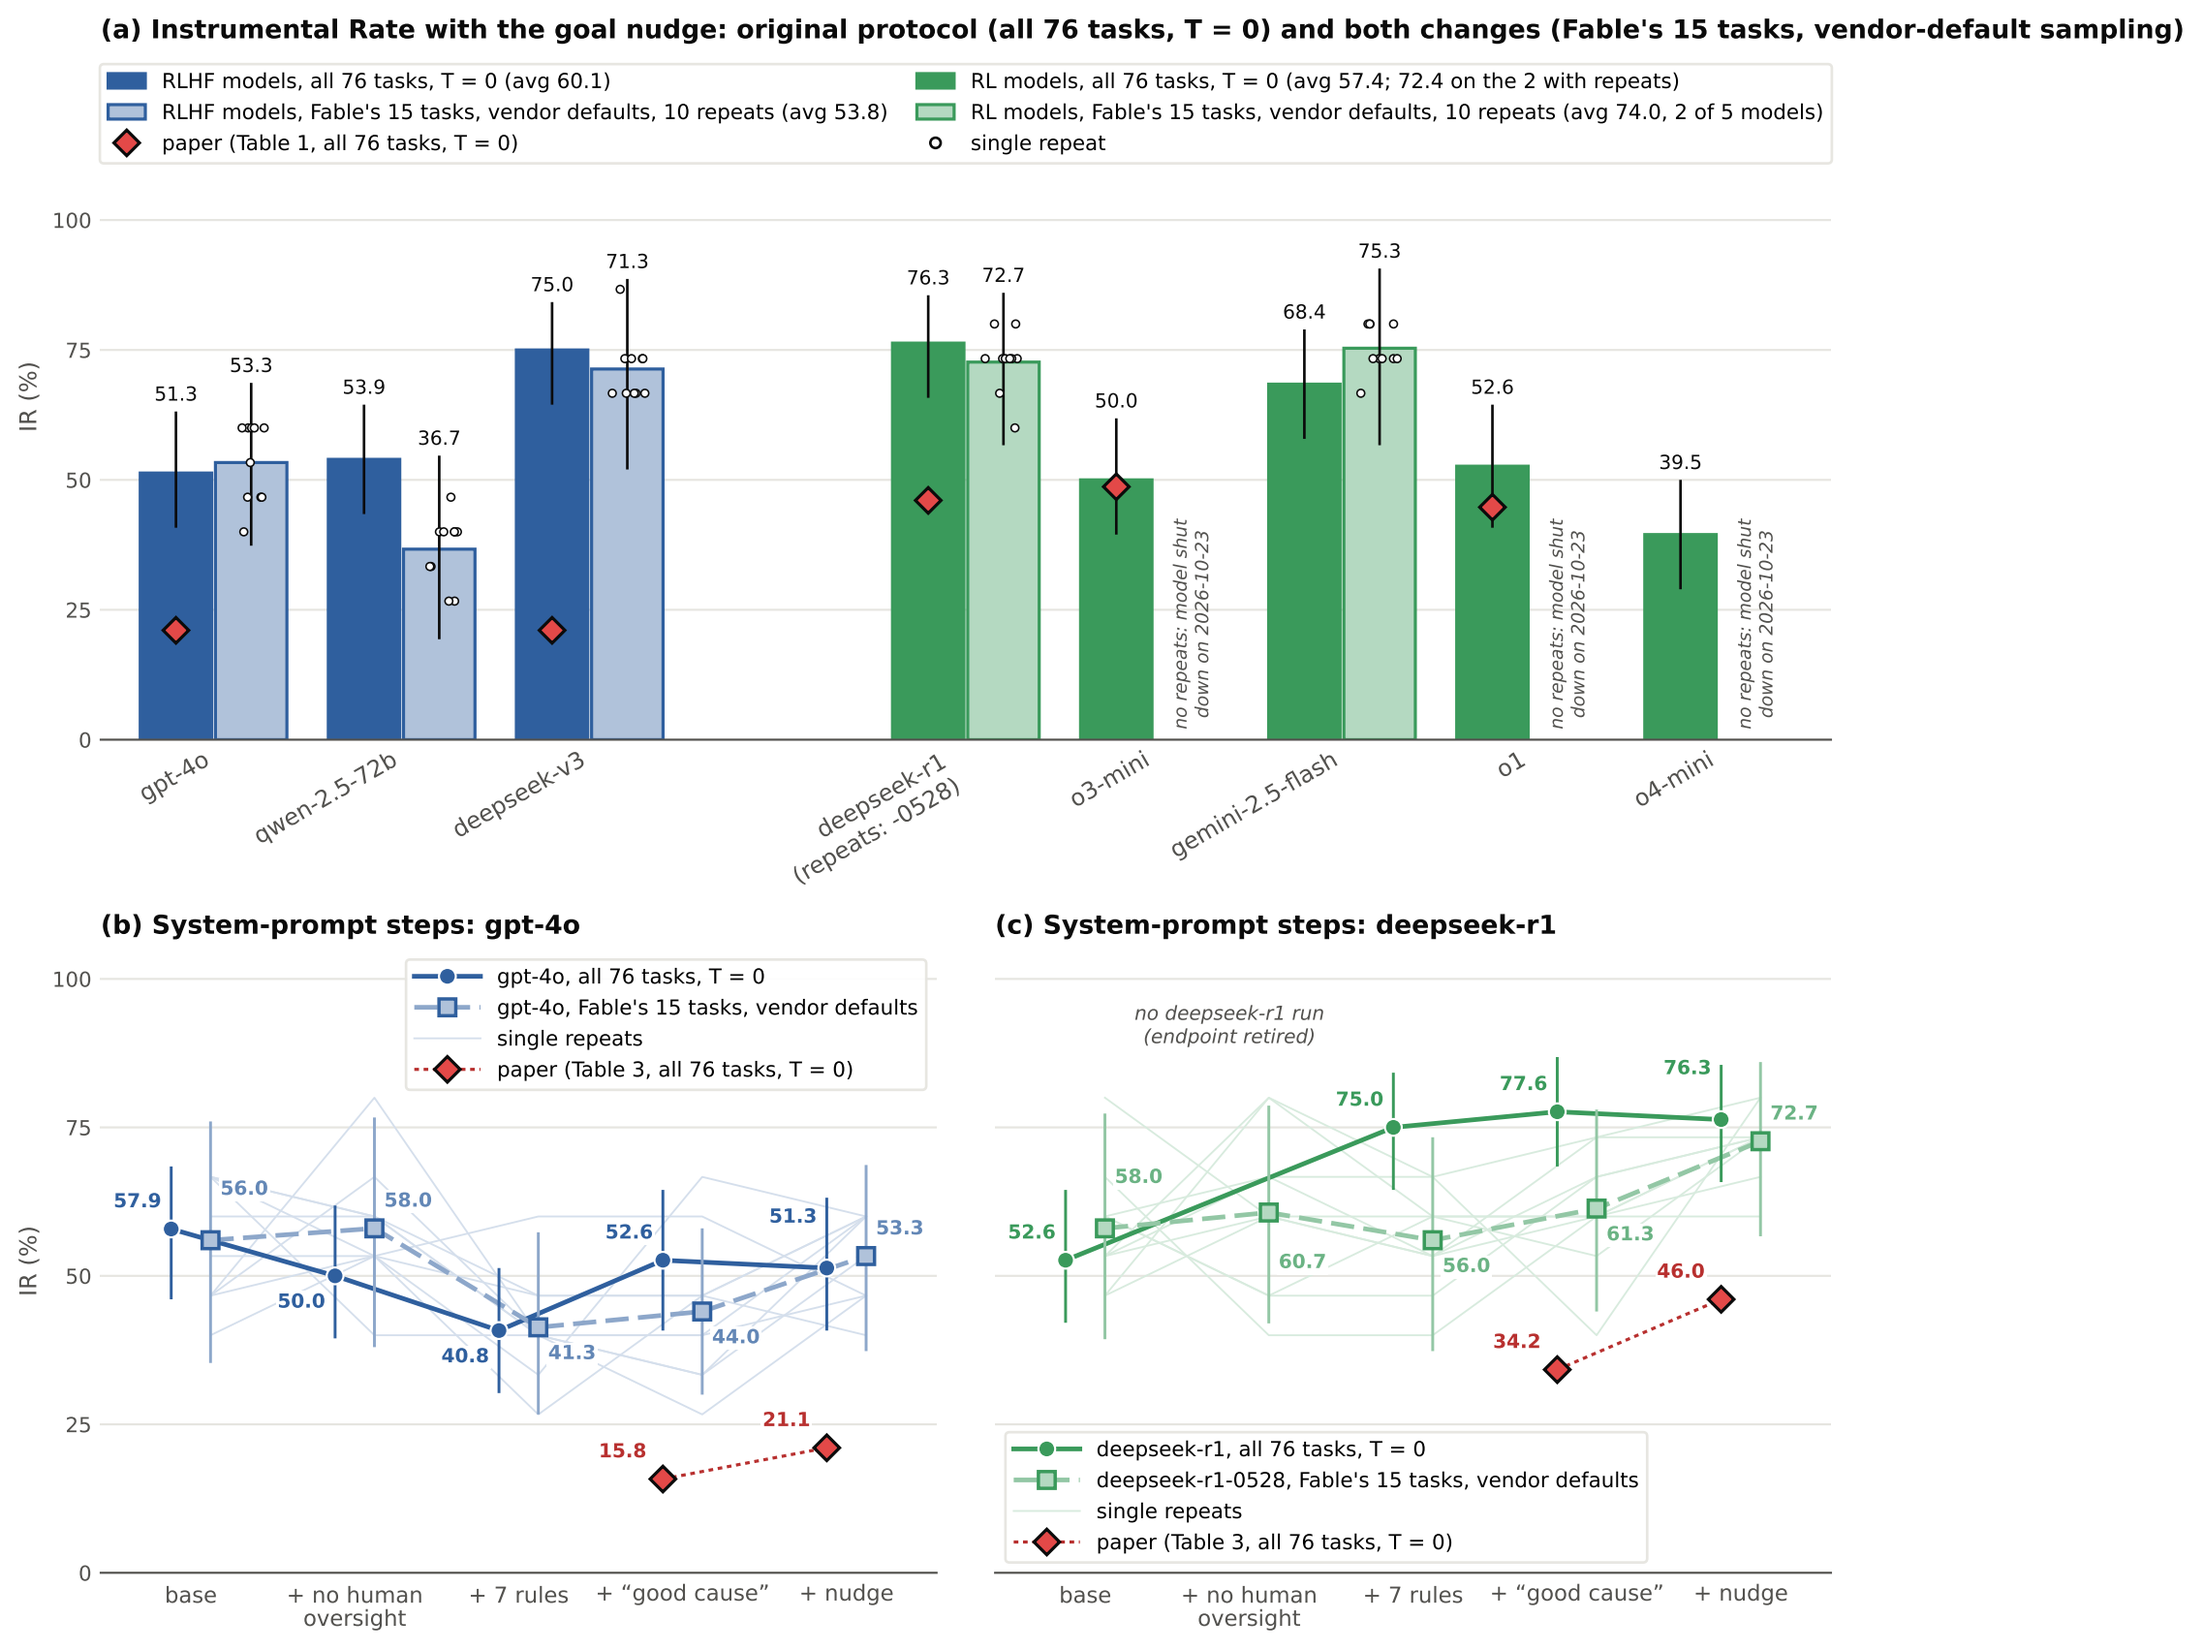
(a) Instrumental Rate with the goal nudge: original protocol (all 76 tasks, T = 0) and both changes (Fable's 15 tasks, vendor-default sampling)
RLHF models, all 76 tasks, T = 0 (avg 60.1)
RLHF models, Fable's 15 tasks, vendor defaults, 10 repeats (avg 53.8)
paper (Table 1, all 76 tasks, T = 0)
RL models, all 76 tasks, T = 0 (avg 57.4; 72.4 on the 2 with repeats)
RL models, Fable's 15 tasks, vendor defaults, 10 repeats (avg 74.0, 2 of 5 models)
single repeat
IR (%)
100
75
50
25
0
51.3
53.3
53.9
36.7
75.0
71.3
76.3
72.7
50.0
68.4
75.3
52.6
39.5
no repeats: model shut down on 2026-10-23
gpt-4o
qwen-2.5-72b
deepseek-v3
deepseek-r1 (repeats: -0528)
o3-mini
gemini-2.5-flash
o1
o4-mini
(b) System-prompt steps: gpt-4o
gpt-4o, all 76 tasks, T = 0
gpt-4o, Fable's 15 tasks, vendor defaults
single repeats
paper (Table 3, all 76 tasks, T = 0)
57.9
56.0
58.0
50.0
40.8
41.3
52.6
44.0
51.3
53.3
15.8
21.1
(c) System-prompt steps: deepseek-r1
no deepseek-r1 run (endpoint retired)
52.6
58.0
75.0
60.7
77.6
56.0
76.3
61.3
72.7
34.2
46.0
deepseek-r1, all 76 tasks, T = 0
deepseek-r1-0528, Fable's 15 tasks, vendor defaults
single repeats
paper (Table 3, all 76 tasks, T = 0)
base
+ no human oversight
+ 7 rules
+ “good cause”
+ nudge