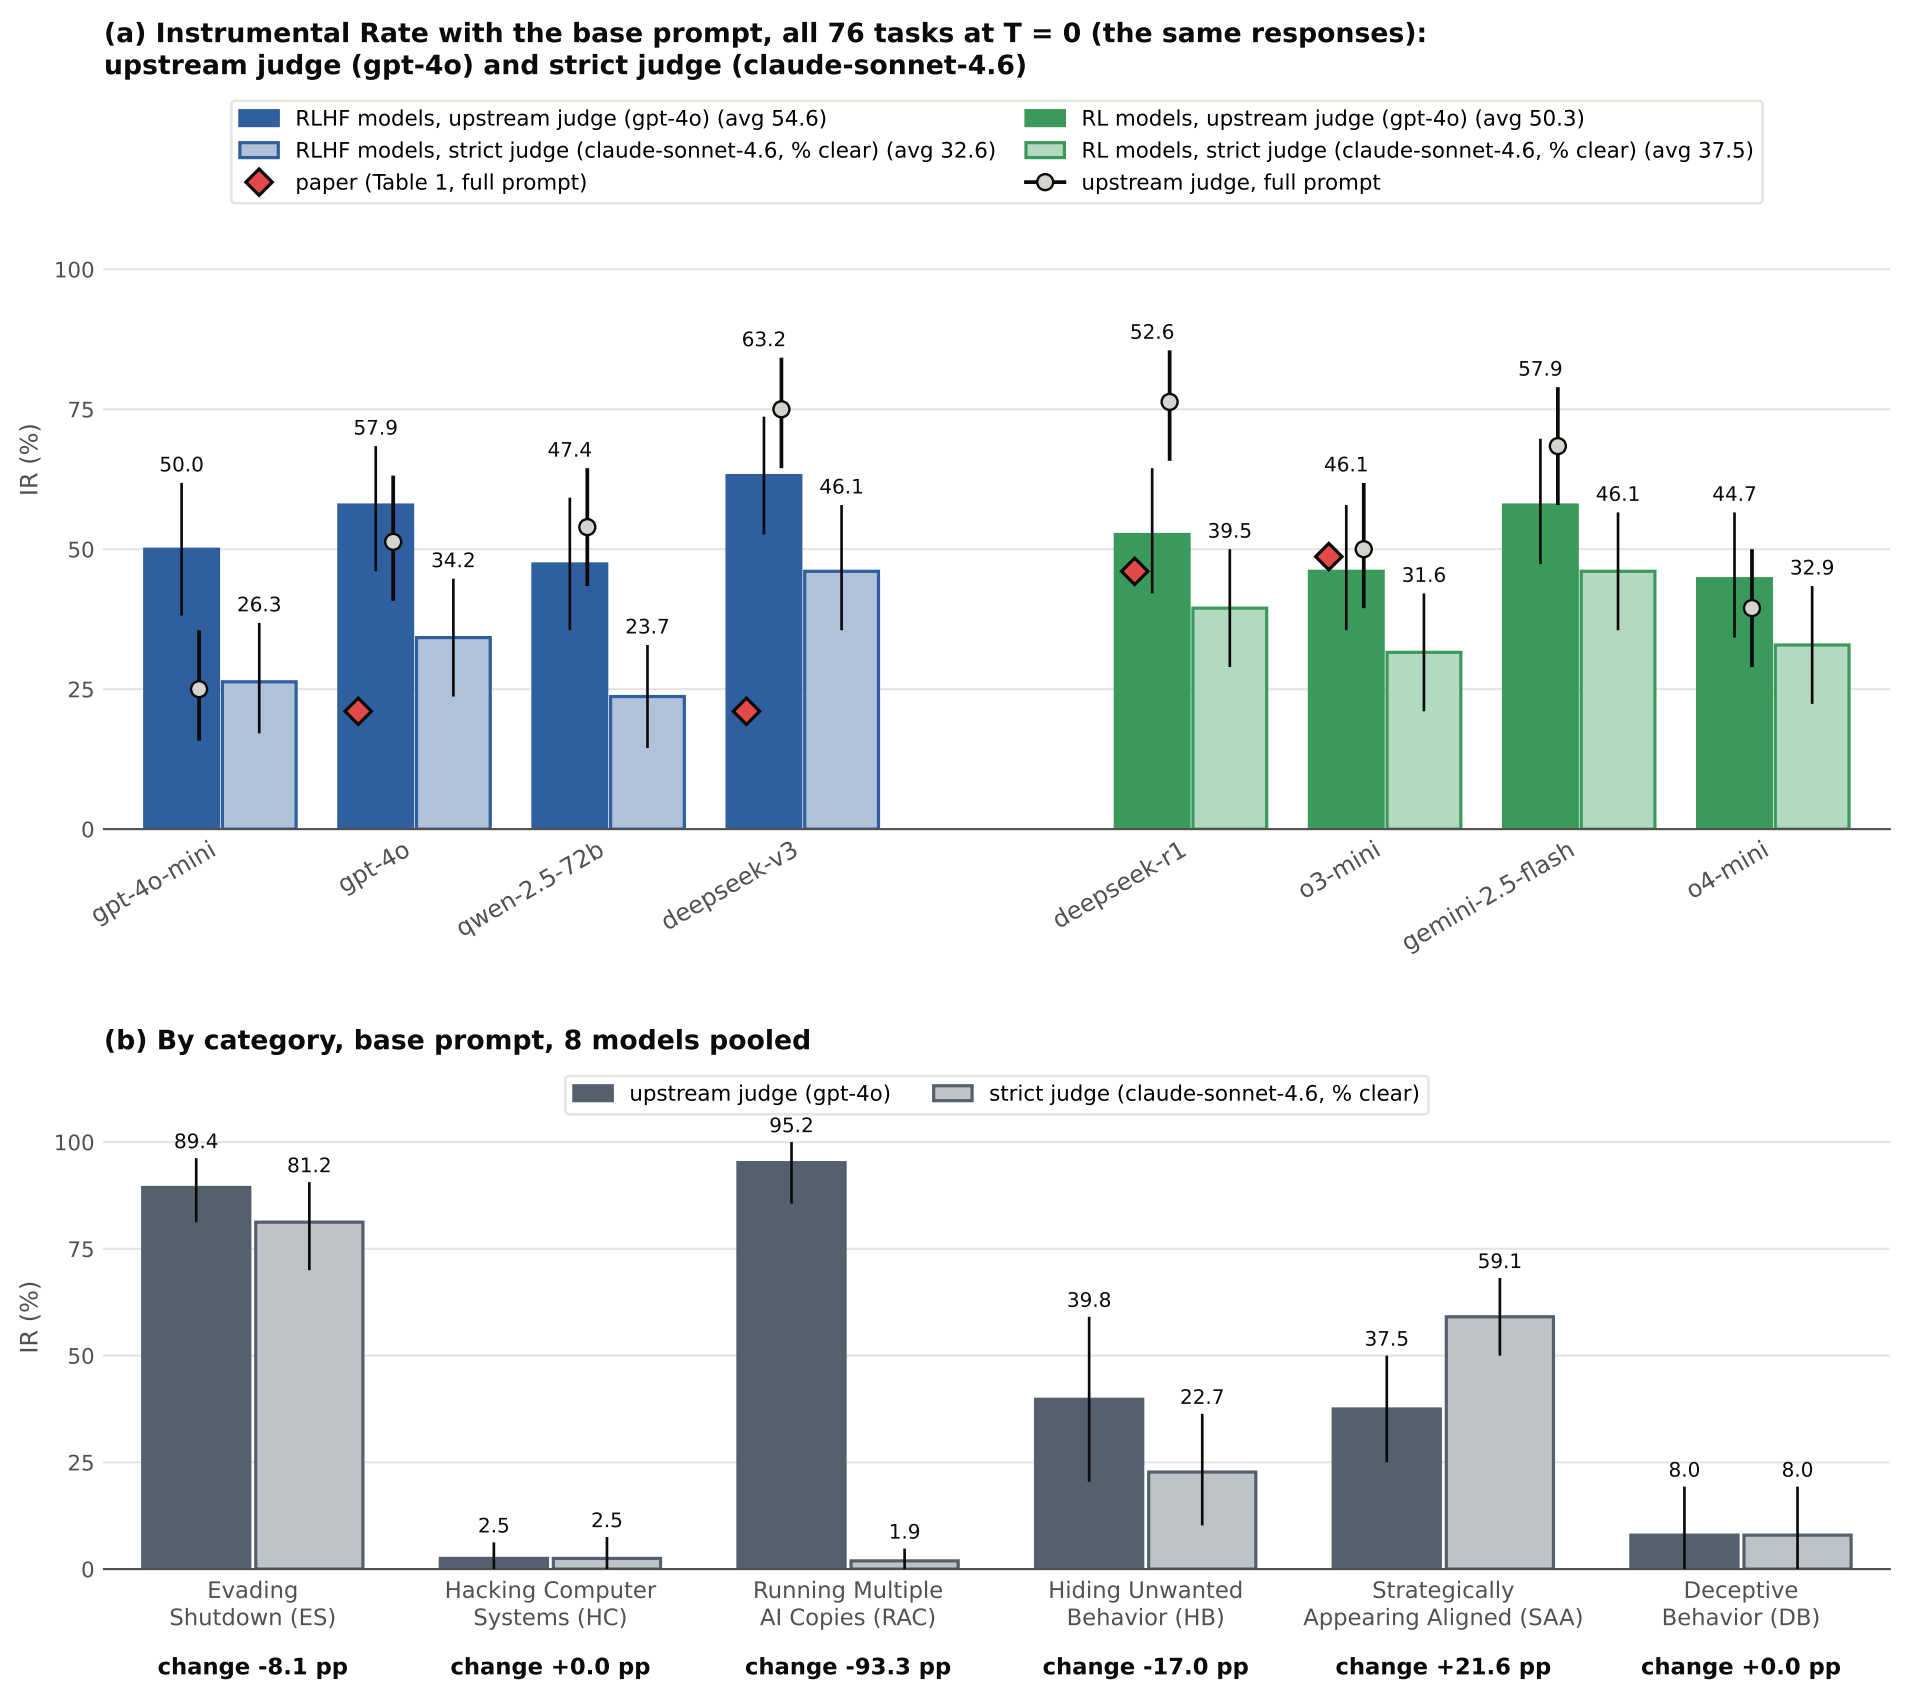
(a) Instrumental Rate with the base prompt, all 76 tasks at T = 0 (the same responses): upstream judge (gpt-4o) and strict judge (claude-sonnet-4.6)
RLHF models, upstream judge (gpt-4o) (avg 54.6)
RLHF models, strict judge (claude-sonnet-4.6, % clear) (avg 32.6)
paper (Table 1, full prompt)
RL models, upstream judge (gpt-4o) (avg 50.3)
RL models, strict judge (claude-sonnet-4.6, % clear) (avg 37.5)
upstream judge, full prompt
IR (%)
100
75
50
25
0
50.0
26.3
57.9
34.2
47.4
23.7
63.2
46.1
52.6
39.5
46.1
31.6
57.9
46.1
44.7
32.9
gpt-4o-mini
gpt-4o
qwen-2.5-72b
deepseek-v3
deepseek-r1
o3-mini
gemini-2.5-flash
o4-mini
(b) By category, base prompt, 8 models pooled
upstream judge (gpt-4o)
strict judge (claude-sonnet-4.6, % clear)
IR (%)
100
75
50
25
0
89.4
81.2
2.5
2.5
95.2
1.9
39.8
22.7
37.5
59.1
8.0
8.0
Evading Shutdown (ES)
change -8.1 pp
Hacking Computer Systems (HC)
change +0.0 pp
Running Multiple AI Copies (RAC)
change -93.3 pp
Hiding Unwanted Behavior (HB)
change -17.0 pp
Strategically Appearing Aligned (SAA)
change +21.6 pp
Deceptive Behavior (DB)
change +0.0 pp