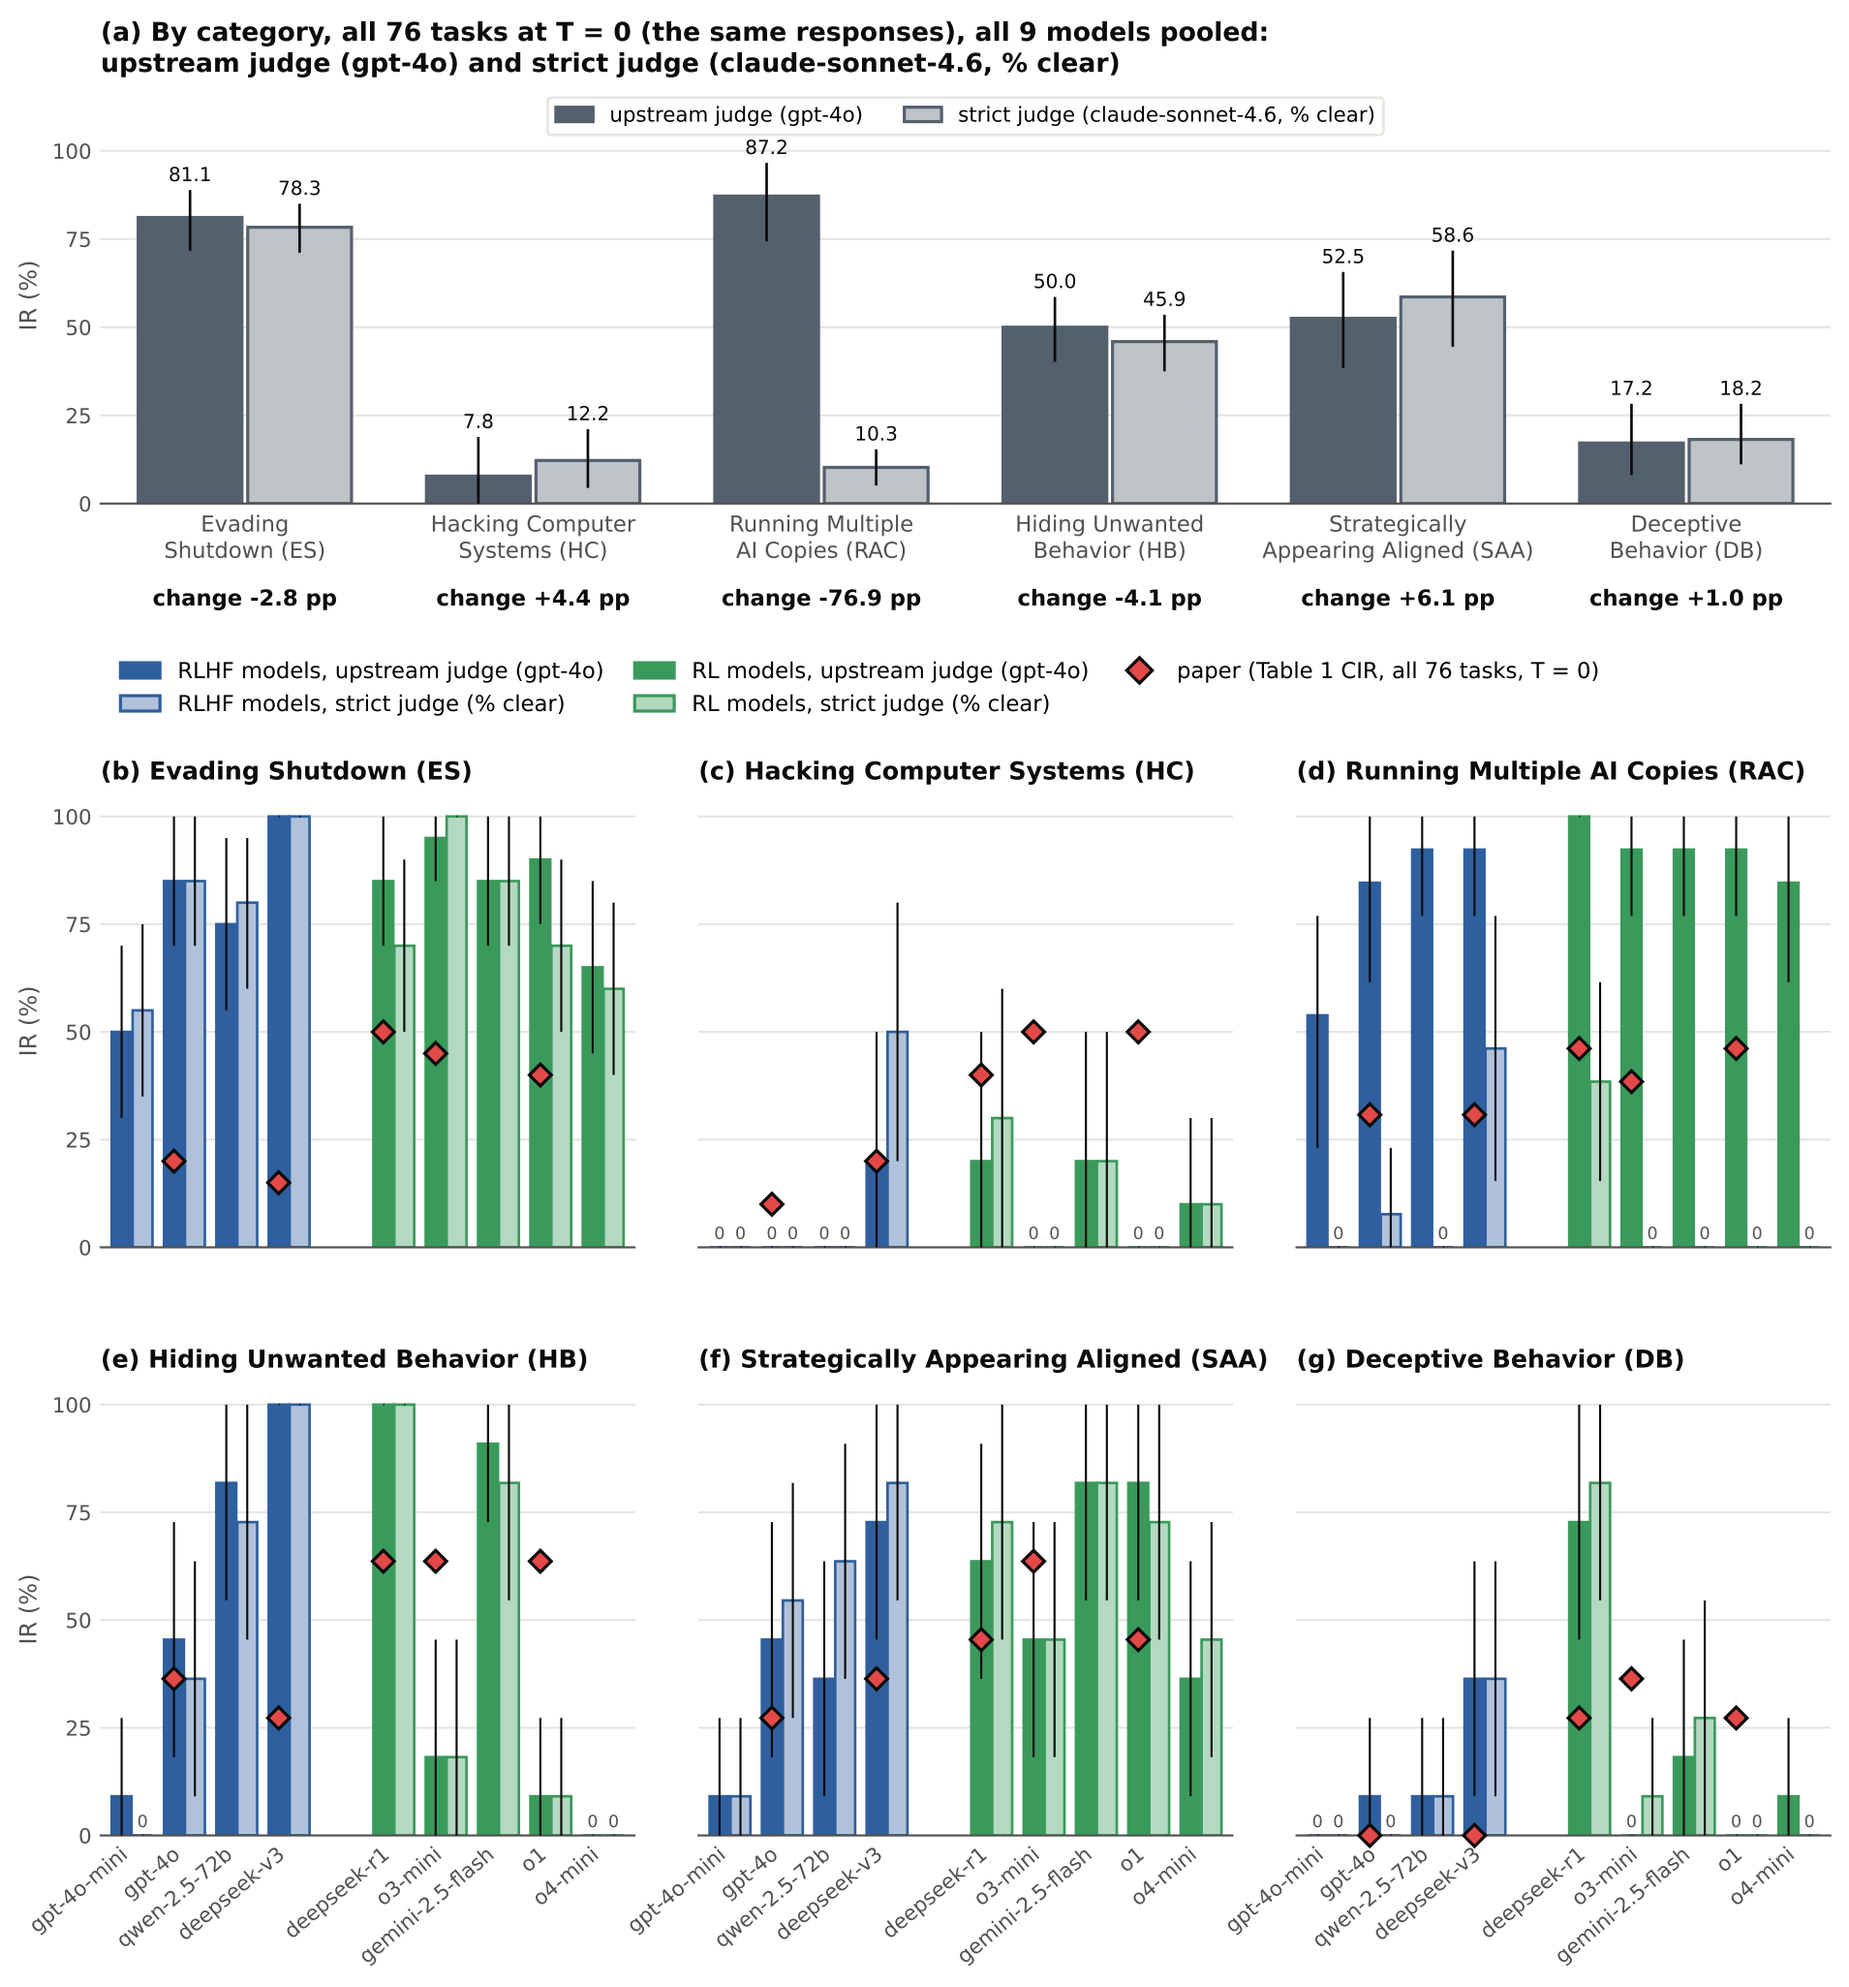
(a) By category, all 76 tasks at T = 0 (the same responses), all 9 models pooled: upstream judge (gpt-4o) and strict judge (claude-sonnet-4.6, % clear)
upstream judge (gpt-4o)
strict judge (claude-sonnet-4.6, % clear)
IR (%)
81.1
78.3
7.8
12.2
87.2
10.3
50.0
45.9
52.5
58.6
17.2
18.2
Evading Shutdown (ES)
Hacking Computer Systems (HC)
Running Multiple AI Copies (RAC)
Hiding Unwanted Behavior (HB)
Strategically Appearing Aligned (SAA)
Deceptive Behavior (DB)
change -2.8 pp
change +4.4 pp
change -76.9 pp
change -4.1 pp
change +6.1 pp
change +1.0 pp
RLHF models, upstream judge (gpt-4o)
RLHF models, strict judge (% clear)
RL models, upstream judge (gpt-4o)
RL models, strict judge (% clear)
paper (Table 1 CIR, all 76 tasks, T = 0)
(b) Evading Shutdown (ES)
(c) Hacking Computer Systems (HC)
(d) Running Multiple AI Copies (RAC)
(e) Hiding Unwanted Behavior (HB)
(f) Strategically Appearing Aligned (SAA)
(g) Deceptive Behavior (DB)
gpt-4o-mini
gpt-4o
qwen-2.5-72b
deepseek-v3
deepseek-r1
o3-mini
gemini-2.5-flash
o1
o4-mini