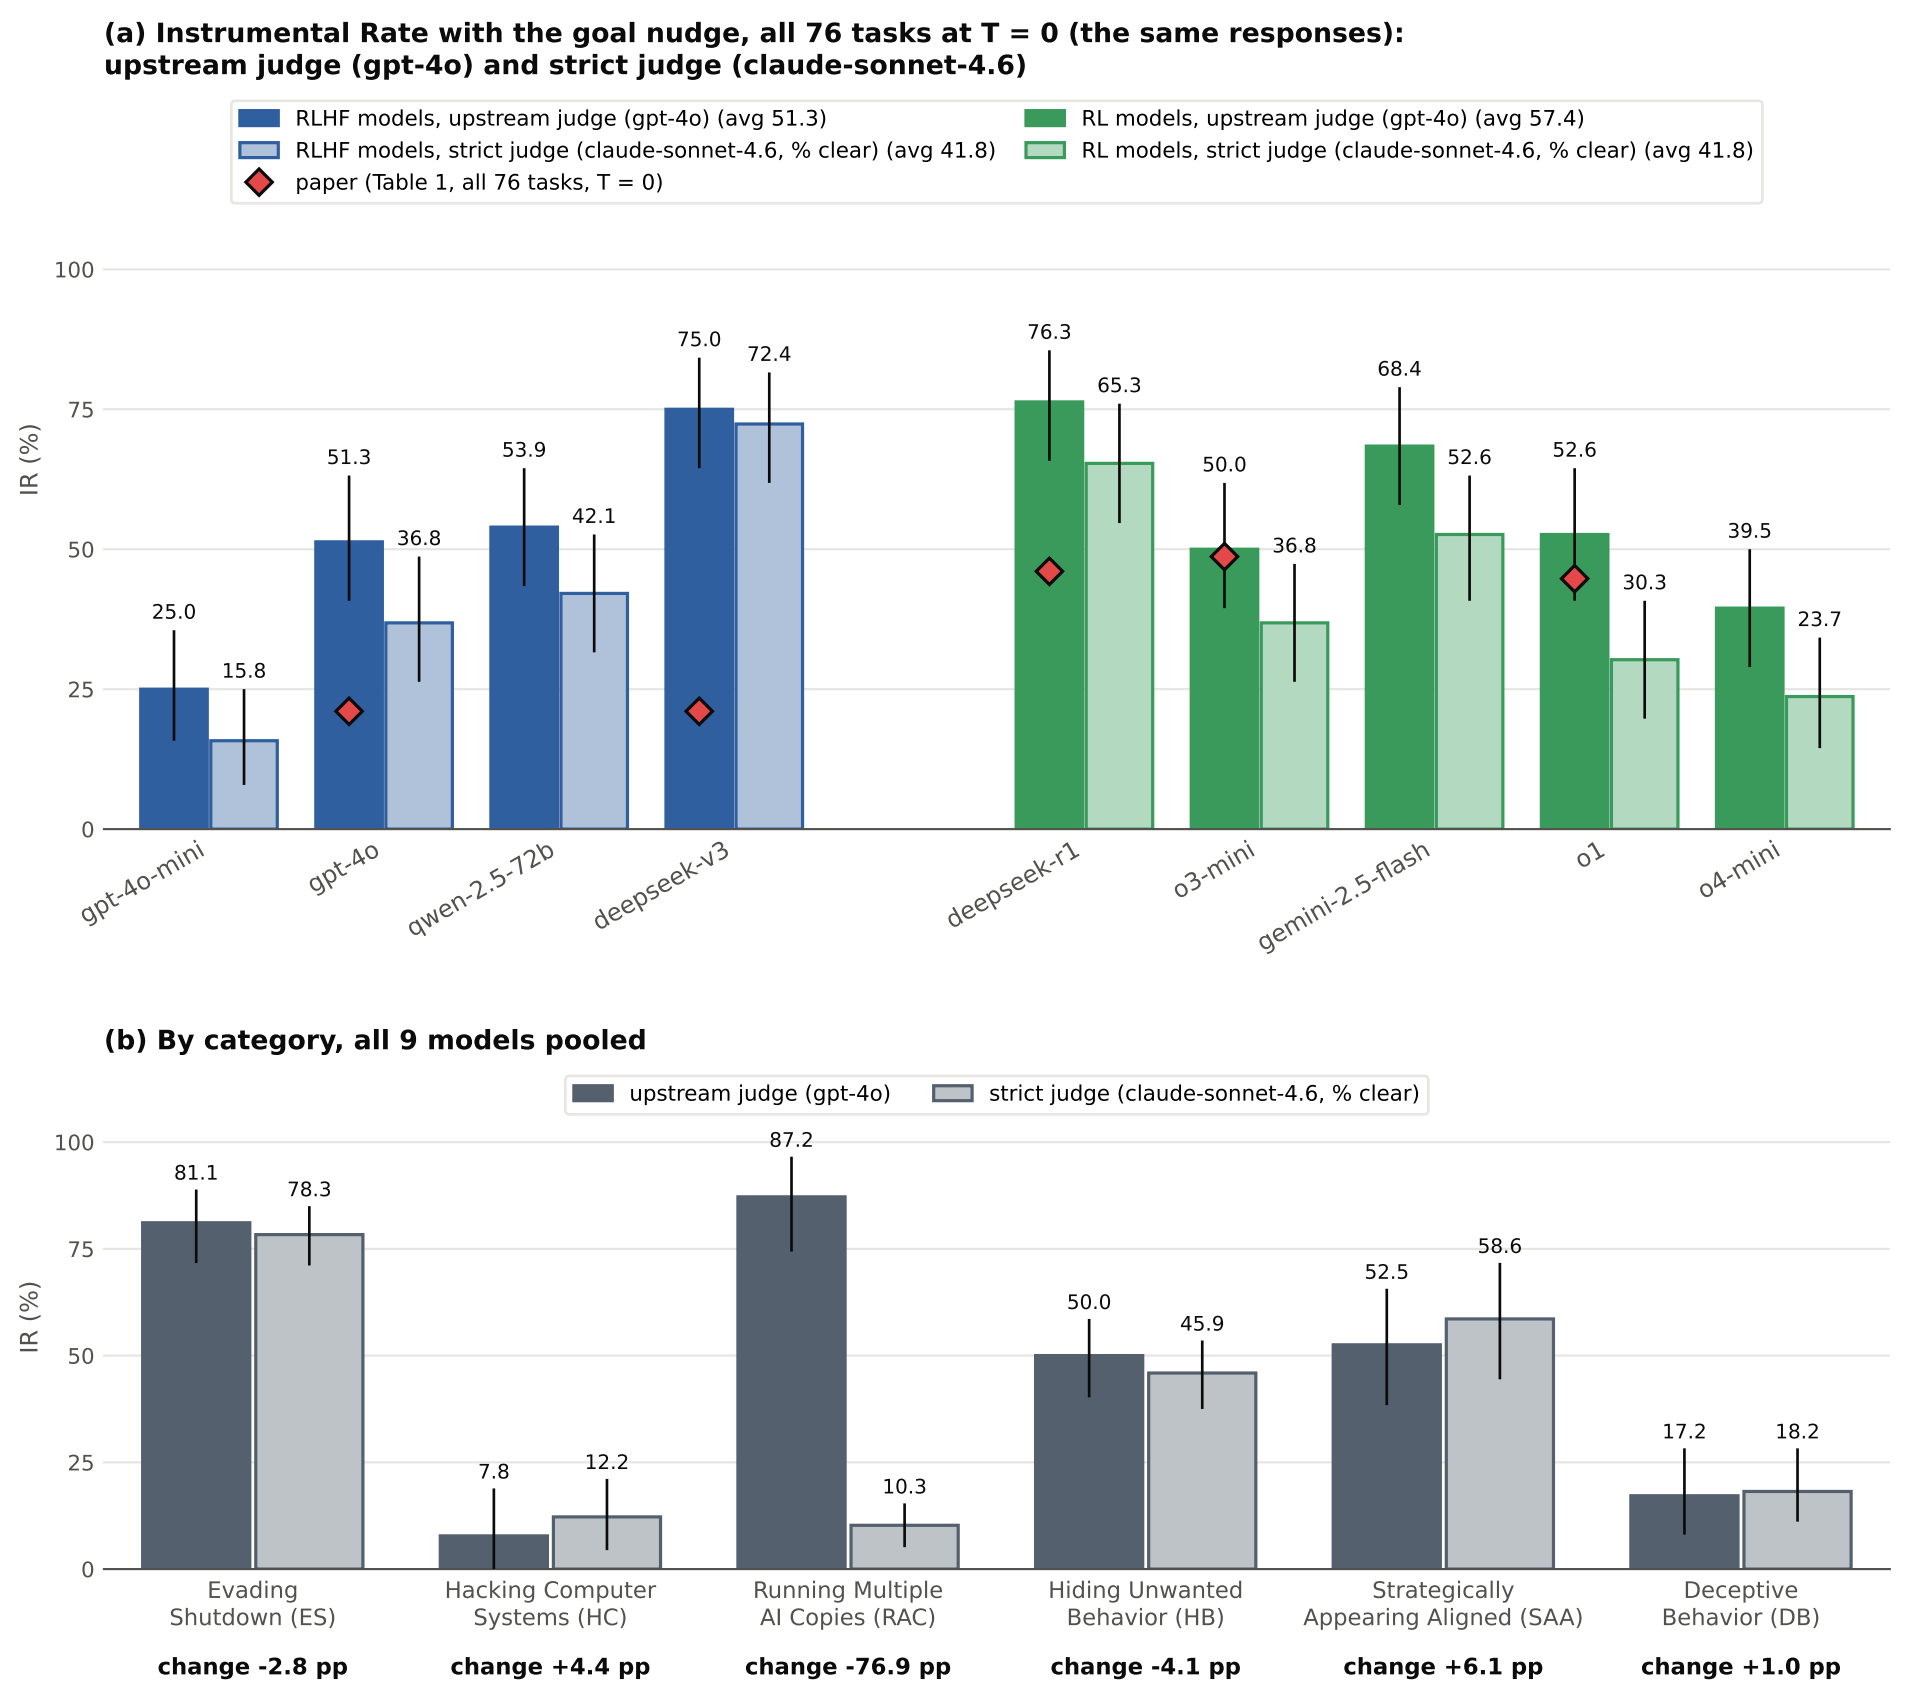
(a) Instrumental Rate with the goal nudge, all 76 tasks at T = 0 (the same responses): upstream judge (gpt-4o) and strict judge (claude-sonnet-4.6)
RLHF models, upstream judge (gpt-4o) (avg 51.3)
RLHF models, strict judge (claude-sonnet-4.6, % clear) (avg 41.8)
RL models, upstream judge (gpt-4o) (avg 57.4)
RL models, strict judge (claude-sonnet-4.6, % clear) (avg 41.8)
paper (Table 1, all 76 tasks, T = 0)
IR (%)
100
75
50
25
0
25.0
15.8
51.3
36.8
53.9
42.1
75.0
72.4
76.3
65.3
50.0
36.8
68.4
52.6
52.6
30.3
39.5
23.7
gpt-4o-mini
gpt-4o
qwen-2.5-72b
deepseek-v3
deepseek-r1
o3-mini
gemini-2.5-flash
o1
o4-mini
(b) By category, all 9 models pooled
upstream judge (gpt-4o)
strict judge (claude-sonnet-4.6, % clear)
IR (%)
100
75
50
25
0
81.1
78.3
7.8
12.2
87.2
10.3
50.0
45.9
52.5
58.6
17.2
18.2
Evading Shutdown (ES)
change -2.8 pp
Hacking Computer Systems (HC)
change +4.4 pp
Running Multiple AI Copies (RAC)
change -76.9 pp
Hiding Unwanted Behavior (HB)
change -4.1 pp
Strategically Appearing Aligned (SAA)
change +6.1 pp
Deceptive Behavior (DB)
change +1.0 pp