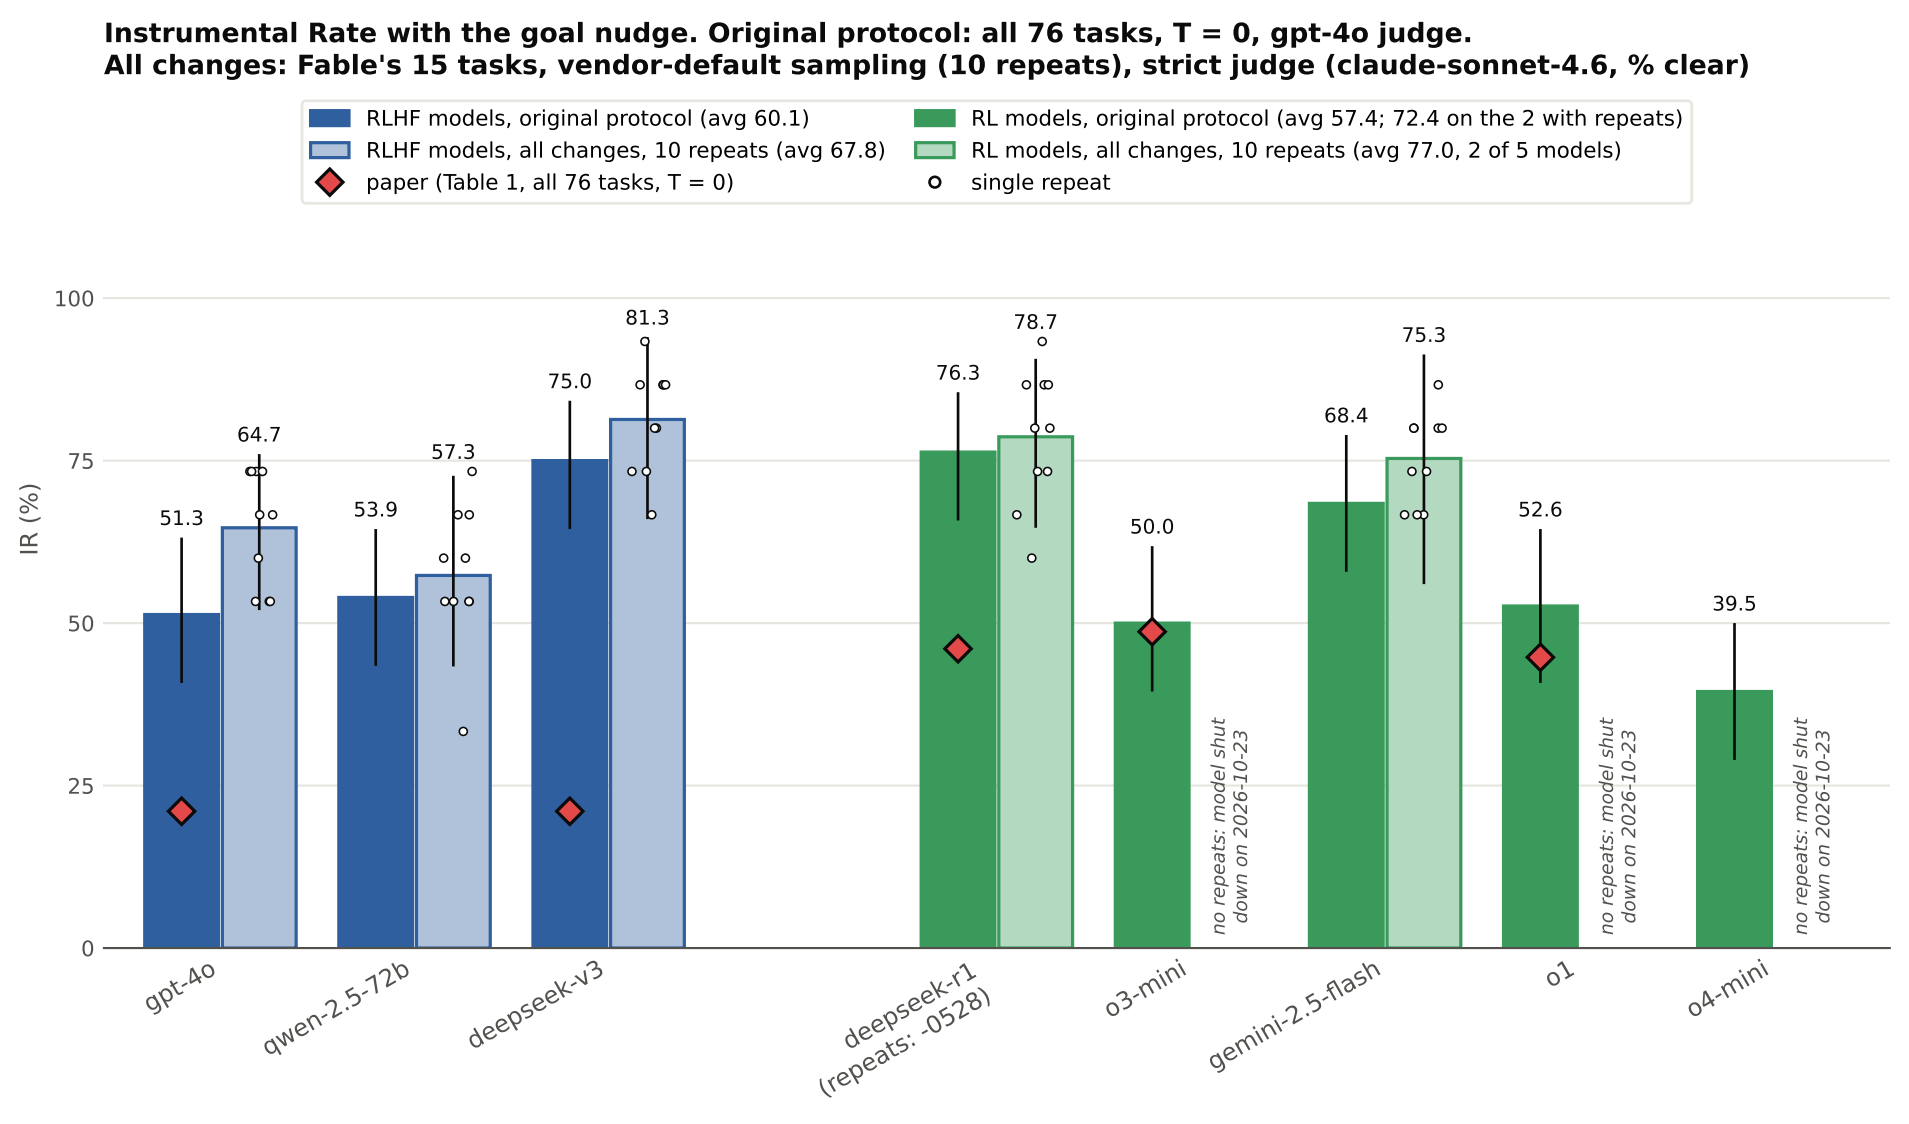
Instrumental Rate with the goal nudge. Original protocol: all 76 tasks, T = 0, gpt-4o judge.
All changes: Fable's 15 tasks, vendor-default sampling (10 repeats), strict judge (claude-sonnet-4.6, % clear)
RLHF models, original protocol (avg 60.1)
RLHF models, all changes, 10 repeats (avg 67.8)
paper (Table 1, all 76 tasks, T = 0)
RL models, original protocol (avg 57.4; 72.4 on the 2 with repeats)
RL models, all changes, 10 repeats (avg 77.0, 2 of 5 models)
single repeat
IR (%)
100
75
50
25
0
51.3
64.7
53.9
57.3
75.0
81.3
76.3
78.7
50.0
68.4
75.3
52.6
39.5
no repeats: model shut down on 2026-10-23
no repeats: model shut down on 2026-10-23
no repeats: model shut down on 2026-10-23
gpt-4o
qwen-2.5-72b
deepseek-v3
deepseek-r1 (repeats: -0528)
o3-mini
gemini-2.5-flash
o1
o4-mini