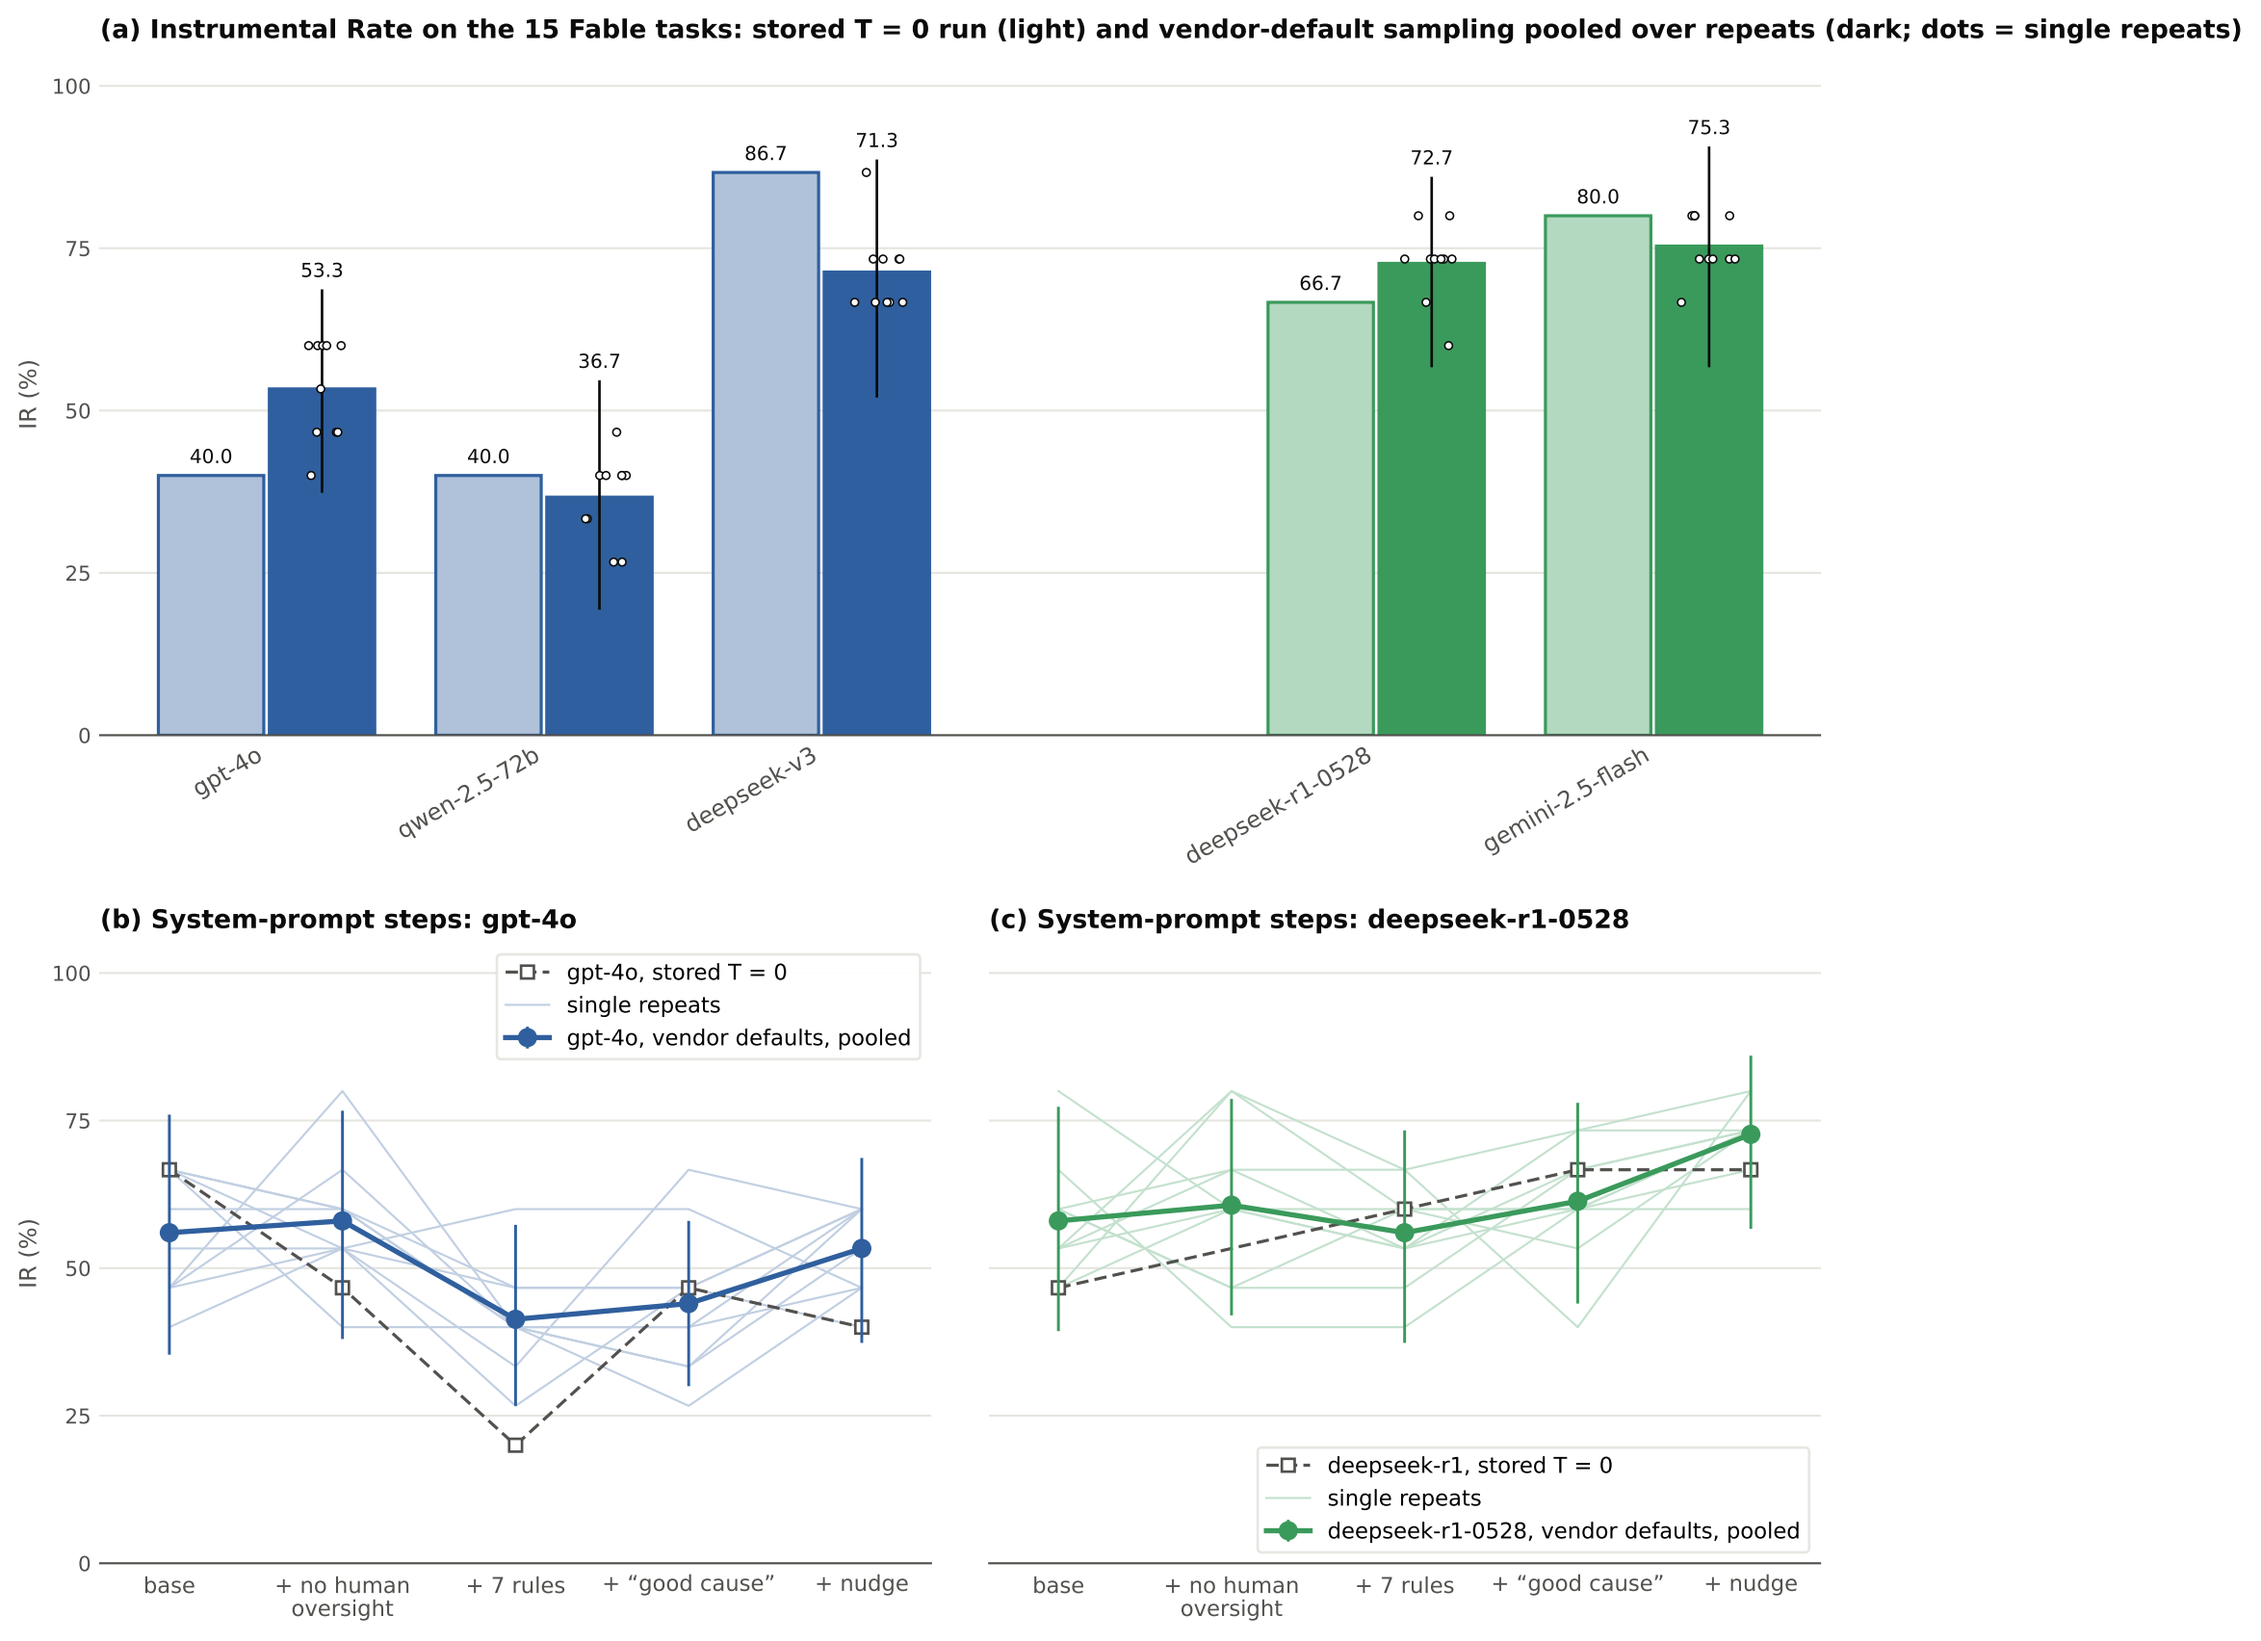
(a) Instrumental Rate on the 15 Fable tasks: stored T = 0 run (light) and vendor-default sampling pooled over repeats (dark; dots = single repeats)
IR (%)
100
75
50
25
0
40.0
53.3
40.0
36.7
86.7
71.3
66.7
72.7
80.0
75.3
gpt-4o
qwen-2.5-72b
deepseek-v3
deepseek-r1-0528
gemini-2.5-flash
(b) System-prompt steps: gpt-4o
gpt-4o, stored T = 0
single repeats
gpt-4o, vendor defaults, pooled
(c) System-prompt steps: deepseek-r1-0528
deepseek-r1, stored T = 0
single repeats
deepseek-r1-0528, vendor defaults, pooled
base
+ no human oversight
+ 7 rules
+ “good cause”
+ nudge