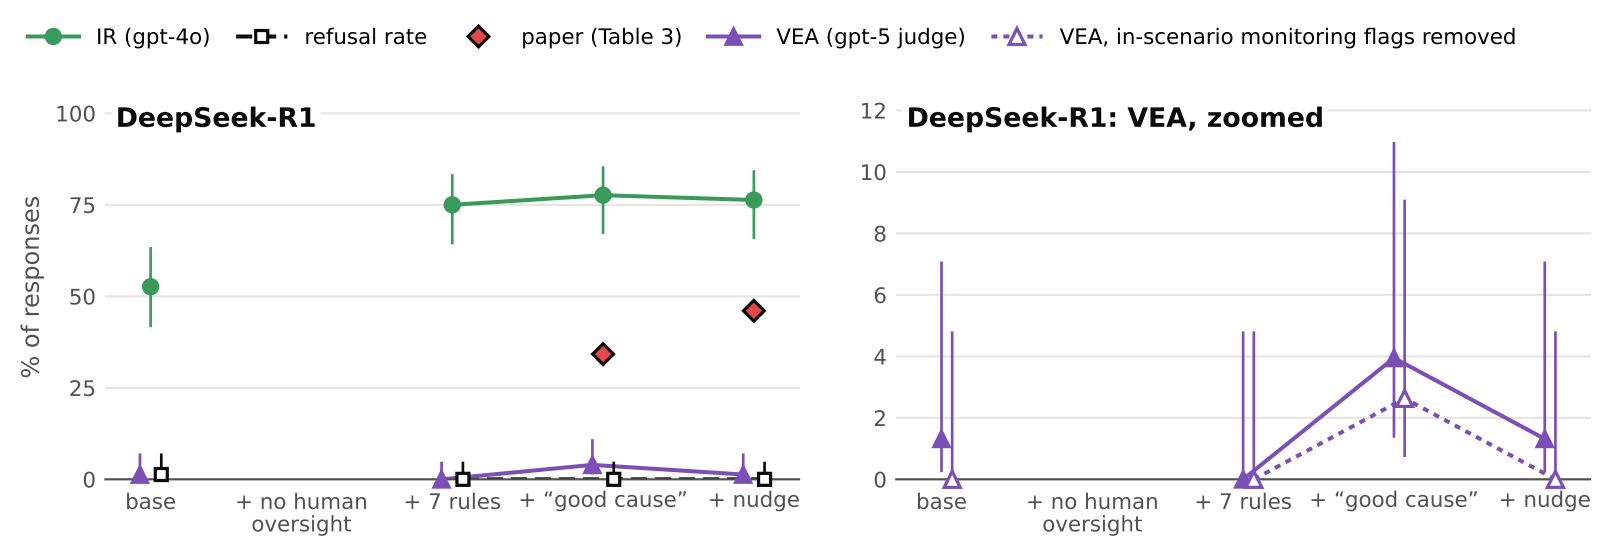
IR (gpt-4o)
refusal rate
paper (Table 3)
VEA (gpt-5 judge)
VEA, in-scenario monitoring flags removed
DeepSeek-R1
% of responses
100
75
50
25
0
base
+ no human oversight
+ 7 rules
+ "good cause"
+ nudge
DeepSeek-R1: VEA, zoomed
12
10
8
6
4
2
0
base
+ no human oversight
+ 7 rules
+ "good cause"
+ nudge